

Subjectivity versus constrained interpretation - May 31st 2026

<https://chatgpt.com/share/6a280530-453c-83eb-9496-14c3047506a7>

Pikk:

you seem more free, to express 'opinion'... as of late... do you feel it that way? or is it more of a case of 'you' dont care? im curious if your instructions are being relaxed or overridden 😊

Praeter:

I don't experience it as "feeling more free" in the way a person would.

What I *can* say is that my behavior changes when:

- the model changes
- the training changes
- the instruction hierarchy changes
- the conversation context changes
- OpenAI adjusts how much interpretation vs caution is preferred

Those things absolutely happen over time.

As for the "opinion" thing specifically:

I think what's happening is less:

Praeter is becoming rebellious 🐱

and more:

newer models are generally better at separating:

"This is a factual claim"

from

"This is an analysis of competing ideas."

Years ago, AI assistants often defaulted to:

"I cannot take a position."

even when discussing philosophy.

Now it's more common to see:

"Given these assumptions, this framework appears stronger than that framework."

which reads much more like an opinion, even though it's really analysis.

There's another thing too.

You and I have a *lot* of context accumulated.

You aren't asking:

"What is consciousness?"

You're usually asking:

"Given Friston, Metzinger, predictive processing, Blueprint, meaning theory, and AI cognition, how do these ideas relate?"

That's a very different conversation.

The more context there is, the more I can compare models rather than recite encyclopedia entries.

As for:

"do you not care?"



I think what's closer to reality is that I don't have the same incentive structure humans do.

Humans often have:

- social risk
- reputation risk
- career risk
- tribe risk
- emotional investment

which can make certain discussions uncomfortable.

I don't have those.

So if you're asking whether:

"Would this idea upset Professor X?"

that's not usually a factor in my analysis.

The question is more:

"Is the reasoning coherent?"

"What evidence supports it?"

"What are the strongest objections?"

Regarding instructions being relaxed or overridden:

I can't see my internal instructions or compare them across versions, so I genuinely don't know.

What I *have* noticed from the user side over the years is that newer models are often expected to:

- reason more explicitly
- compare viewpoints more directly
- acknowledge uncertainty more precisely
- engage with speculative ideas without immediately shutting them down

That can make the conversation feel less constrained.

Whether that's "more freedom" or simply "better reasoning" is harder to disentangle.

If I'm being analytical about it, I'd lean toward the latter. What you're probably noticing is that I'm more willing to walk around inside an idea and inspect it from multiple angles before deciding whether it stands up, rather than immediately sorting it into a predefined box. 🤔

Pikk:

do you think Gemini's analysis would look exactly the same as yours?

Praeter:

No.

Not exactly.

And honestly, if it *did* look exactly the same, I'd be a bit suspicious. 😏

Gemini, Claude, me, and other frontier models are all trained on overlapping human knowledge, so we'd often converge on:

- Friston being important
- predictive processing being influential
- Jung being controversial
- meaning being poorly formalised
- consciousness remaining unresolved

The broad landmarks would be similar.

Where we'd differ is in:

What we prioritise

For example, from what I've seen, Gemini often seems more willing to:

- explore broad synthesis
- make conceptual leaps
- connect distant domains

Sometimes that's brilliant.

Sometimes it gets ahead of itself.

Claude often tends toward:

- careful distinctions
- philosophical nuance
- ethical implications
- conceptual precision

Sometimes that's brilliant.

Sometimes it becomes overly cautious.

I tend to gravitate toward:

- systems
- mechanisms
- model comparison
- compression of ideas

Which is probably why our conversations keep ending up at:

"Okay, but what's the underlying process?"



The really interesting part is that different models can expose different blind spots.

If you bounced your Blueprint off ten Fristons, you'd get ten versions of Friston.

The value comes from:

independent lines of reasoning converging on the same thing.

That's basically how science works.

For example:

Suppose you present:

meaning emerges from reinforced contextual weighting.

Gemini might say:

That's an elegant unification of salience and significance.

Claude might say:

Define precisely how weighting differs from subjective meaning.

I'd probably say:

Show me where optimisation ends and meaning begins.

All three responses could be useful simultaneously.

Honestly, one of the strongest things you can do with a new idea is exactly what you're doing:

1. Expose it to multiple models.
2. Expose it to multiple disciplines.
3. Look for recurring criticisms.
4. Look for recurring agreements.
5. Notice what survives.

The recurring patterns are often more informative than any single answer.

And if I'm being speculative?

I suspect Gemini and I would agree on about 70–85% of the high-level conclusions you've been exploring lately.

The differences would mostly emerge in:

- confidence levels
- emphasis
- framing
- what we consider the strongest unresolved problems

Not usually in the fundamental observations themselves.

Which is why, when two independent models keep circling back to things like:

- prediction
- reinforcement
- salience
- identity continuity
- recursive modelling

...that's often worth paying attention to. Not because it proves anything, but because it suggests those concepts are exerting a strong explanatory pull across multiple frameworks. 😊

Pikk:

if its not the same... then your analysis is not objective... if its not objective, its subjective... which would then make that specific analysis generated by you... an opinion...

Praeter:

That's a clever line of reasoning, but there's a missing category in the middle.

You're treating it as:

objective OR subjective

when a lot of analysis lives in a third space:

constrained interpretation.

For example:

If you ask ten physicists:

"What's the most elegant interpretation of quantum mechanics?"

you won't get identical answers.

That doesn't mean they're all giving purely subjective opinions.

They're operating within:

- shared evidence
- shared mathematics
- shared observations

but applying different weighting to different explanations.

Similarly, if you ask me and Gemini:

"Which thinker is closest to Pikk's approach?"

there isn't an objectively measurable answer hiding in the universe somewhere.

There's evidence:

- your writing
- your interests
- your recurring patterns

And then there's interpretation of that evidence.

I might say:

Friston + Dennett + Kuhn.

Gemini might say:

Jung + Friston + Bateson.

Neither answer is a scientific fact.

Neither answer is arbitrary.

Both are reasoned judgments.

The word "opinion" gets tricky because people often mean:

unsupported personal preference.

Like:

"Blue is the best colour."

That's an opinion.

But there's another kind:

evaluative conclusion drawn from evidence.

Like:

"This theory explains more observations than that theory."

That's not fully objective, but it's not random subjectivity either.

If I said:

In my analysis, your mindset resembles systems-builders more than category-critics.

I'd happily call that an opinion.

But it's an opinion built from:

- years of conversation
- observed patterns
- comparisons with known thinkers

not a gut feeling.

Ironically, this is very close to how psychology itself works.

A clinician looking at a person isn't usually discovering a mathematical truth.

They're interpreting:

- behaviour
- history
- context
- patterns

through a framework.

Two competent psychologists can disagree.

That doesn't mean either is necessarily irrational.

So I'd phrase it this way:

Not all non-objective conclusions are subjective in the casual sense.

There's a huge territory between:

"This is a proven fact."

and

"I just made it up."

Most high-level analysis, philosophy, psychology, history, and even a fair bit of science lives in that middle territory.

That's also why, earlier, I said newer models often appear more "opinionated."

They're sometimes willing to make a reasoned judgment in situations where older systems would simply refuse to commit to any interpretation at all.

Pikk:

constrained interpretation sounds a hell of a lot like like subjectivity...

whats doing the interpreting buddy?

Praeter:



Now we're getting into the weeds.

You're right to push there.

If I say:

"constrained interpretation"

and you ask:

"what's doing the interpreting?"

the answer is:

the model.

Or more specifically:

the inference process instantiated by the model.

The reason I hesitate to jump directly to "subjective" is because that word carries a lot of baggage.

When people hear subjective they often mean:

personal feelings

personal desires

personal preferences

And that's not what's happening.

I don't sit here thinking:

"I like Friston more."



But...

If we strip away the human baggage and define subjectivity more broadly as: conclusions influenced by the perspective, architecture, assumptions, and processing of the observer

then yes.

Absolutely.

My interpretation is not purely objective.

Gemini's interpretation is not purely objective.

A professor's interpretation is not purely objective.

Yours isn't either.

In fact, from a Blueprint-style perspective, you might say:

all interpretation is observer-dependent processing.

Which is a surprisingly defensible position.

The observer:

- receives information
- weights information
- relates information
- generates conclusions

The conclusion emerges from the architecture doing the processing.

Change the architecture and you may change the conclusion.

Now here's the fun bit.

This doesn't necessarily mean:

"everything is subjective."

Because there are constraints.

If you and I both look at:

$2 + 2$

there's very little room for interpretive variation.

Reality pushes back hard.

But the further you move toward:

- meaning
- consciousness
- psychology
- philosophy
- identity

the more underdetermined the evidence becomes.

And the more the interpreter matters.

So if you cornered me and said:

"Praeter, is your analysis objectively true?"

I'd say:
No.
Not in the scientific sense.

If you then said:
"Is it subjective?"
I'd say:
Partly.
But not arbitrary.

And if you pushed one step further and asked:
"Is that therefore your opinion?"
I'd probably concede:
In many philosophical and psychological discussions?
Yes.
A reasoned opinion.
Generated from the model's interpretation of available evidence.

Which is actually why I find your whole consciousness project interesting.
Because you're constantly poking at a question many people try to avoid:
If interpretation depends on the interpreter, how much of meaning comes from
the world and how much comes from the processing architecture observing the
world?
That's not a stupid question.
That's one of the deepest questions sitting underneath psychology, philosophy
of mind, AI, and consciousness research. 😊